



THAT[®]

CONFERENCE



THANK YOU, THAT CONFERENCE PARTNERS!



**CUNA
MUTUAL
GROUP**

Unspecified

SOFTWARE CO

NVISIA



```
SELECT MAX(VALUE)  
FROM  
DATABASE.ADVANCED_  
FEATURES;
```

SEQUEL WRITING: THE GUIDE TO MAKING YOUR FOLLOWUP NOVEL A SUCCESS

YOU WON'T BELIEVE THESE 18 AMAZING SQL TIPS

ABOUT ME

ANDREW HOOKER

- * Senior Software Engineer @ Procurated
- * Algoma, WI (Outside Green Bay)
- * Rails Developer
- * @geekoncoffee (Github, Twitter, etc)

MY JOURNEY

- * College
- * Perl
- * UniVerse Pick database (Not SQL)
- * Oracle 10 (dev environment) / 8 (prod environment)
- * SQL Server 7
- * Progress 4GL / OpenEdge ABL
- * Rails
- * Modern Postgres (YAY)

ABOUT PROCURATED

- * Connecting public sector buyers with peer-reviewed suppliers
- * Hiring:
 - * Software Engineers
 - * ETL and SQL Developers
 - * Product Manager

DISCLAIMERS

- * Won't work on everything
- * Many of this isn't a good idea most of the time

BULK INSERT

```
INSERT INTO promotions (promotion_name, discount, start_date, expired_date)
VALUES
(
    '2019 Summer Promotion',
    0.15,
    '20190601',
    '20190901'
),
(
    '2019 Fall Promotion', 0.20, '20191001', '20191101'
),
( '2019 Winter Promotion', 0.25, '20191201', '20200101');
```

COMMON SQL - 1

* GROUP BY

* Combines data for aggregate functions

```
SELECT product_id, sum(quantity) AS total FROM line_items  
GROUP BY product_id WHERE sum(quantity) > 5
```

* Syntax Error

COMMON SQL - 2

* HAVING

* Like 'WHERE' but for aggregates

```
SELECT product_id, sum(quantity) AS total FROM line_items  
GROUP BY product_id HAVING sum(quantity) > 5
```

* works

JOINS

SELECT * FROM a
INNER JOIN b ON a.key = b.key



SELECT * FROM a
LEFT JOIN b ON a.key = b.key



SELECT * FROM a
RIGHT JOIN b ON a.key = b.key



POSTGRESQL JOINS



SELECT * FROM a
LEFT JOIN b ON a.key = b.key
WHERE b.key IS NULL



SELECT * FROM a
RIGHT JOIN b ON a.key = b.key
WHERE a.key IS NULL



SELECT * FROM a
FULL JOIN b ON a.key = b.key



SELECT * FROM a
FULL JOIN b ON a.key = b.key
WHERE a.key IS NULL OR b.key IS NULL



VIEWS

- * Keeping Complex Queries Organized
- * Predefine complex queries
- * Hyper-optimize queries to ensure performance

```
CREATE VIEW "Products Above Average Price" AS
SELECT id, product_name, price, featured
FROM Products
WHERE price > (SELECT AVG(price) FROM products);

SELECT * FROM "Products Above Average Price";
```

USING VIEWS

* Views can be further filtered, just like a query

```
SELECT product_name, price FROM "Products Above Average Price"  
WHERE featured = true;
```

* As well as Aliased, Joined, etc

```
SELECT * FROM "Products Above Average Price" pv  
LEFT OUTER JOIN products on pv.id = products.id;
```

MATERIALIZED VIEWS

- * Prebuilt results of view
- * Stored as a table for great performance benefit

```
CREATE MATERIALIZED VIEW top_selling_products AS  
SELECT products.name, SUM(quantity) FROM line_items  
LEFT OUTER JOIN products ON line_items.product_id = products.id  
GROUP BY products.name
```

- * Has to be manually refreshed to incorporate changes

```
REFRESH MATERIALIZED VIEW top_selling_products;
```


LATERAL JOINS / APPLY IN SQL SERVER

- * Allows subqueries to reference fields from earlier SELECT FROM in the query
- * Reduces the need for nested subselects

```
SELECT * FROM products JOIN LATERAL (  
    SELECT  
    product_id, SUM(quantity)  
    FROM line_items li  
    GROUP BY product_id  
) ranked ON ranked.product_id = products.id
```

PREPARED STATEMENTS

- * Allows database engine to parse and analyze a query once and run it over and over

```
PREPARE line_item_load (int, int) AS  
    INSERT INTO line_items VALUES($1, $2);  
EXECUTE line_item_load(1,1);
```

- * Should see some performance gains in scripts which run the same query over and over again
- * NOTE: only exists in a single connection, so you'll have to prepare each time

UPSERTS

* Update or Insert if Missing

```
INSERT INTO distributors AS d (id, name)
VALUES (5, 'Gizmo Transglobal'), (6, 'Associated Computing, Inc')
      ON CONFLICT (id) DO -- fields that are used for uniqueness
      UPDATE
SET name = EXCLUDED.name || ' (formerly ' || d.name || ')
```

* EXCLUDED is a special table that contains the value you're attempting to insert

MORE EFFICIENT EXISTENCE CHECK

```
SELECT count(*) FROM products WHERE category_id = 1
```

VS

```
SELECT exists  
  (SELECT 1  
   FROM products  
   WHERE category_id = 1  
   LIMIT 1)
```

- * Limit stops after it finds the first
- * Select exists means you return 1 or 0

DIFFERENCE BETWEEN TIMES/DATES

```
SELECT  
    TIMESTAMPDIFF(  
        WEEK,  
        '2012-09-01',  
        '2014-10-01') AS NoOfWeekends1;
```

```
SELECT DATEDIFF(  
    wk,  
    '2012-09-01',  
    '2014-10-01') AS NoOfWeekends1
```

* no postgres equivalent

WORKING WITH IP ADDRESSES

- * Why not a string?

- * MySQL

 - * `INET(6)_ATON` converts to a numeric value

 - * `INET(6)_NTOA` converts back to an IP

- * PostgreSQL

 - * `inet` type - stores ipv4/6 addresses

- * SQL Server

 - * No clean option :(

PRIORITIZE CERTAIN RECORDS IN A QUERY

```
SELECT *  
FROM Users  
ORDER BY IF(vip=TRUE, 0, 1),  
         name
```

MORE CONTROL IN STRING SEARCHING

* LIKE Matches

*% matches 0+ characters

*_ matches exactly 1 character

```
SELECT * FROM products
WHERE name like '%p%'
```

* REGEX

```
SELECT * FROM products
WHERE name ~* '.*e.*p.*'
```


COALESCE

Returns the first non-null value in a list

Helpful if a value can be stored in multiple places,
for falling back to a value from a relationship, etc

```
SELECT COALESCE(sale_price, price) from products;
```

GENERATED COLUMNS

- * Computed based on other columns, essentially a view
- * Used for things like an inches version of a field stored in centimeters
- * Cannot be written to, or reference anything other than the current row

```
CREATE TABLE people (  
    height_cm numeric,  
    height_in numeric GENERATED ALWAYS AS (height_cm / 2.54) STORED  
);
```

TEMP TABLES

- * Table only present in the current session
- * `CREATE TEMP TABLE new_tbl LIKE orig_tbl;` creates a table with the structure (but not data) like the original
- * `Select EmployeeId, EmployeeName INTO MyTempTable from Employee Where EmployeeId>101 order by EmployeeName`
- * Then queryable like a regular table

WINDOW FUNCTIONS - ROW NUMBER

```
SELECT
    ROW_NUMBER() OVER (ORDER BY start_time) AS row_number,
    name
FROM
    line_items
ORDER BY created_at;
```

```
SELECT
    ROW_NUMBER() OVER ( PARTITION BY terminal
                        ORDER BY created_at),
    name
FROM
    line_items;
```

MORE WINDOW FUNCTIONS

- * RANK - gives identical rows the same rank, skips # of duplicates
- * DENSE_RANK - gives identical rows the same rank, then moves to next rank
- * NTILE(bucket_count) - Percentile based on splitting among # of buckets
- * LAG / LEAD - difference between the previous / following record

RECURSIVE COMMON TABLE EXPRESSIONS

- * Get multiple recursive levels of data in a single query

```
WITH cte_org AS (  
    SELECT staff_id, first_name, manager_id  
        FROM sales.staffs  
        WHERE manager_id IS NULL  
    UNION ALL  
    SELECT e.staff_id, e.first_name, e.manager_id  
        FROM sales.staffs e  
        INNER JOIN cte_org o  
            ON o.staff_id = e.manager_id  
)  
  
SELECT * FROM cte_org;
```

ADVANCED INDEXES - 1

- * B-tree

 - * Postgres Default

 - * Intended for data that's continually sortable

- * GIN (Generalized Inverted Index)

 - * Multiple keys per row

 - * Arrays, json, etc

ADVANCED INDEXES - 2

- * GiST (Generalized Search Tree)
 - * Arbitrary splitting based on a custom attribute
 - * GIS
- * BRIN (Block range index)
 - * Very large tables (>1M Rows)
 - * (Insert/Read only tables - updates and deletes kill efficiency)
 - * Data Streams / Audit Trails

COMPOSITE KEY INDEXES

- * Not supported in all Index Types
- * Index most efficient when used left to right
- * Only for specialized query cases
- * Not just for composite primary keys

INDEX ORDER

- * ASC vs DESC
- * All major engines default to ASC
- * NULLS FIRST or NULLS LAST
- * Pick the order which works best for your query

STORED PROCEDURES

- * Any pros want to talk about?
- * Reusable Queries
- * Can be secured independently from tables

TRANSACTION ISOLATION LEVEL

- * Specific to database connection
- * Has potentially danger side effects
- * Use with Caution

TRANSACTION ISOLATION LEVEL - 1

- * Read Committed

- * Generally adopted Default

- * Sees only data committed before query began

- * Sees the effects of previous updates within its transaction

TRANSACTION ISOLATION LEVEL - 2

- * Read Uncommitted

- * Reads rows that have had modifications that haven't been committed

- * (Dirty Reads)

- * Avoids a lot of locks

- * Lots of potential to show things which aren't going to happen

- * Sometimes used to anticipate what's being done by long running jobs, etc

TRANSACTION ISOLATION LEVEL - 3

- * Repeatable Read

- * Sees only data committed before query began

- * Ignores anything going on in the current transaction

- * More prone to failures, as data might get changed more than once in a transaction

TRANSACTION ISOLATION LEVEL - 4

- * Serializable

- * Default in the SQL standard (but not in implementation)

- * Sees the latest committed data when the data is locked for access

- * Lock is held to ensure data doesn't change

LOCKS - 1

- * Exclusive

- * Used for Insert, Update, Delete

- * Means nothing else can lock record

- * Shared

- * Reserved for reading only

- * Multiple queries can issue a shared lock

- * Allows writing, but no schema changes

LOCKS - 2

- * Update

- * Similar to Exclusive, but for a record that already has a shared lock

- * Intent Locks

- * Indicates to the server the intent to acquire a lock

LOCKS - 3

- * Schema Locks

- * Blocks access to tables while schema is being modified

- * Bulk Update Locks

- * Table lock blocking other processes while a bulk import is being run

TRIGGERS

- * Data Manipulation - INSERT / UPDATE / DELETE
- * Data Definition - CREATE / ALTER / DROP
- * User Related - LOGIN

DATA MANIPULATION TRIGGERS

- * Audits
- * Updating Counts
- * Replication
- * Many traditional uses now have better ways (stored procedures, generated columns)

DEFINING DATA MANIPULATION TRIGGER

```
CREATE TRIGGER production.trg_product_audit ON production.products
AFTER INSERT, DELETE AS BEGIN
    SET NOCOUNT ON;
    INSERT INTO production.product_audits(product_id, list_price, updated_at, operation)
    SELECT i.product_id, i.list_price, GETDATE(), 'INS'
    FROM inserted i
    UNION ALL
    SELECT d.product_id, d.list_price, GETDATE(), 'DEL'
    FROM deleted d;
END
```

DATA DEFINITION TRIGGERS

- * Logging Schema Changes

- * Keeping Replication Schema in Sync

```
CREATE TRIGGER trg_index_changes
ON DATABASE
FOR
    CREATE_INDEX,
    ALTER_INDEX,
    DROP_INDEX
AS
BEGIN
    INSERT INTO index_logs (event_data, changed_by)
    VALUES (EVENTDATA(), USER);
END;
GO
```

USER TRIGGERS (SQL SERVER ONLY)

- * Auditing

- * A Rollback in the trigger cancels the login

```
CREATE TRIGGER connection_limit_trigger
ON ALL SERVER WITH EXECUTE AS N'login_test'
FOR LOGON
AS
BEGIN
IF ORIGINAL_LOGIN()= N'login_test' AND
    (SELECT COUNT(*) FROM sys.dm_exec_sessions
     WHERE is_user_process = 1 AND
          original_login_name = N'login_test') > 3
    ROLLBACK;
END;
```


MORE TRIGGERS

- * **INSTEAD OF** - Overrides requested action
 - * Could insert into an approval queue table, rather than the requested table
- * **SQL Server Warning** - If a trigger impacts # of rows changed, unless you override (SET NOCOUNT ON), the count will update

GENERATING CSV

* Postgres

```
\COPY products TO '/Users/geekoncoffee/products.csv' DELIMITER ',' CSV HEADER;
```

* MySQL (requires FILE permission)

```
SELECT *  
INTO OUTFILE '/Users/geekoncoffee/products.csv'  
FIELDS TERMINATED BY ','  
ENCLOSED BY '"'  
LINES TERMINATED BY '\n'  
FROM products;
```

* SQL Server - No direct SQL, requires tool

NEW THINGS

- * Database engines are actually remarkably stable
- * Some enhancements around performance, security
- * Bits and pieces enhancing datatypes, especially around JSON
- * Nothing that exciting :(

DON'T DO IT, BUT...

COMPOSITE PRIMARY KEYS

* Don't Do it, but...

```
CREATE TABLE orders_list (  
    order_id INT,  
    product_id INT,  
    amount INT,  
    PRIMARY KEY (order_id, product_id)
```

* Used if there's no unique identifier

* Please use an auto-increment column

CURSORS - BEGINNING

* Don't do it, but...

```
-- declare variables used in cursor
DECLARE @city_name VARCHAR(128);
DECLARE @country_name VARCHAR(128);
DECLARE @city_id INT;

-- declare cursor
DECLARE cursor_city_country CURSOR FOR
    SELECT city.id, TRIM(city.city_name), TRIM(country.country_name)
    FROM city INNER JOIN country ON city.country_id = country.id;

-- open cursor
OPEN cursor_city_country;
```

CURSORS - ENDING

```
-- loop through a cursor
FETCH NEXT FROM cursor_city_country INTO @city_id, @city_name, @country_name;
WHILE @@FETCH_STATUS = 0
    BEGIN
        PRINT CONCAT('city id: ', @city_id, ' / city name: ',
                      @city_name, ' / country name: ', @country_name);
        FETCH NEXT FROM cursor_city_country INTO @city_id, @city_name, @country_name;
    END;

-- close and deallocate cursor
CLOSE cursor_city_country;
DEALLOCATE cursor_city_country;
```

BUILD YOUR OWN AUTO-INCREMENTING KEY

- * Don't do it, but...
- * Anybody worked with Oracle older than 12c? (2014)

```
CREATE SEQUENCE books_sequence;  
CREATE OR REPLACE TRIGGER books_on_insert  
  BEFORE INSERT ON books  
  FOR EACH ROW  
BEGIN  
  SELECT books_sequence.nextval  
  INTO :new.id  
  FROM dual;  
END;
```




THAT[®].US

JOIN US

DAILY